

Audio Deepfake and Tampering Detection: Project Documentation

Institution: Clarkson University Last Updated: November 2025

Table of Contents

Part I: Research Overview

1. Executive Summary
2. Research Objectives
3. Datasets
4. Model Architectures
5. Audio Tampering Methodology
6. Training Methodology
7. Experimental Results
8. Key Findings

Part II: Reproducibility Guide

1. Environment Setup
2. Dataset Preparation
3. Training Procedures
4. Evaluation Procedures
5. GUI Applications
6. Troubleshooting

Part III: References & Code

1. Scientific Citations
 2. Code Availability
-

Part I: Research Overview

1. Executive Summary

This project implements and evaluates two state-of-the-art audio deepfake detection systems on standard benchmarks and novel tampering attacks. The research demonstrates that self-supervised pre-training significantly improves detection robustness across diverse conditions.

Key Results Summary

Model	ASVspoof 2019 LA	ASVspoof 2021 LA	Trans-Splicing Detection
XLS-R + SLS	0.26% EER	2.97% EER	95.45%
AASIST	0.83% EER	48.27% EER	41.72%

Main Contributions

1. Benchmark Reproduction: Successfully reproduced published results for both AASIST (Jung et al., 2022) and XLS-R + SLS (Zhang et al., 2024)
2. Cross-Dataset Evaluation: Demonstrated XLS-R's superior generalization across clean audio (2019), codec-distorted audio (2021), and tampering attacks
3. In-House Tampering Datasets: Created two novel evaluation datasets:
 - Trans-Splicing Dataset: 1,932 files with TTS-generated word insertions
 - Semantic Tampering Dataset: 50 files with NLP-guided audio modifications
4. Interactive GUI Applications: Developed user-friendly interfaces for both models

Major Finding

XLS-R + SLS outperforms AASIST on BOTH clean and codec-distorted datasets:

- ASVspoof 2019 LA: 0.26% vs 0.83% (3.2x improvement)
- ASVspoof 2021 LA: 2.97% vs 48.27% (16.3x improvement)

2. Research Objectives

2.1 Problem Statement

Audio deepfakes pose significant threats to:

- Voice authentication systems
- Media authenticity verification
- Legal evidence integrity
- Personal privacy and security

2.2 Research Goals

1. Reproduce benchmark results for AASIST on ASVspoof 2019 LA
 2. Reproduce paper results for XLS-R + SLS on ASVspoof 2021 LA
 3. Compare model generalization across different datasets
 4. Evaluate detection of novel tampering techniques
 5. Create interactive tools for practical use
-

3. Datasets

3.1 ASVspoof 2019 LA (Logical Access)

Source: Edinburgh DataShare URL: <https://datashare.ed.ac.uk/handle/10283/3336>

Partition	Bonafide	Spoof	Total	Attacks
Train	2,580	22,800	25,380	A01-A06
Dev	2,548	22,296	24,844	A01-A06
Eval	7,355	63,882	71,237	A07-A19

Audio Specifications: FLAC format, 16kHz sample rate, mono channel, clean (no codec distortion)

3.2 ASVspoof 2021 LA (Logical Access)

Source: Zenodo URL: <https://zenodo.org/record/4837263>

Partition	Bonafide	Spoof	Total
Eval	14,816	133,360	148,176

Codec Distribution:

Codec	Percentage	Description
ulaw	15.9%	G.711 mu-law
opus	15.9%	Opus codec
gsm	15.9%	GSM 06.10
alaw	13.1%	G.711 A-law
none	13.1%	No codec (clean)
pstn	13.1%	Simulated PSTN
g722	13.1%	G.722 wideband

Key Challenge: 86.9% of evaluation samples contain codec distortion, making this dataset significantly more challenging than ASVspoof 2019.

3.3 Trans-Splicing Dataset (In-House)

Total Files: 1,932 tampered audio files **Technique:** Word-level replacement using TTS-generated segments

Category	Files	TTS System	Processing
xtts-clean	506	XTTS (Coqui)	Basic normalization
xtts-unclean	508	XTTS (Coqui)	With artifacts
yourtts-clean	536	YourTTS	Basic normalization
yourtts-unclean	382	YourTTS	With artifacts

3.4 Semantic Tampering Dataset (In-House)

Total Files: 50 (9 bonaf de + 41 tampered) Technique: NLP-guided word deletion at phoneme boundaries

4. Model Architectures

4.1 Architecture Comparison

Feature	AASIST	XLS-R + SLS
Architecture	Graph Networks	Attention Wav2Vec 2.0 + SLS Module
Pre-training	None (from scratch)	436K hours self-supervised
Parameters	~ 300K	~ 340M
Input	Raw waveform (64,600 samples)	Raw waveform (64,600 samples)
Feature Extraction	Sinc Convolution	Contrastive Learning
Classifier	Heterogeneous GAT	Layer Selection + FC
Training Epochs	100	4 (early stopped at 2)
ASVspoof 2019 EER	0.83%	0.26% (BEST)
ASVspoof 2021 LA EER	48.27%	2.97%

4.2 AASIST Architecture

Reference: Jung et al., “AASIST: Audio anti-spoofing using integrated spectro-temporal graph attention networks”, IEEE/ACM TASLP 2022

Architecture Components:

Raw Audio (64,600 samples @ 16kHz)

Sinc Convolution Layer (learnable mel-scale filterbank)

Residual Encoder (6 ResNet-style blocks)

Dual Graph Attention

GAT-S (Spectral)	GAT-T (Temporal)
Max-pool over time	Max-pool over frequency

Heterogeneous GAT (cross-domain integration)

Graph Pooling (top-k node selection)

Classification Head [Bonafide, Spoof]

Parameters: ~297,866 trainable parameters

4.3 XLS-R + SLS Architecture

Reference: Zhang et al., “Audio Deepfake Detection with Self-supervised XLS-R and SLS classifier”, ACM MM 2024

Architecture Components:

Raw Audio (64,600 samples @ 16kHz)

XLS-R 300M Backbone (frozen)

Pre-trained on 436K hours

128 languages

24 transformer layers

SLS Module (Sensitive Layer Selection)

Weighted combination of all layers

Learnable layer importance weights


Classification Head [Bonafide, Spoof]

Parameters: ~340M total (mostly frozen pre-trained weights)

5. Audio Tampering Methodology

5.1 Trans-Splicing Pipeline


Trans-splicing creates tampered audio by replacing specific words with TTS-generated alternatives.



Trans-Splicing Pipeline Diagram

Figure 5.1: Trans-splicing pipeline showing word-level replacement with TTS-generated segments.


Pipeline Steps:



5.2 TTS Systems Used

System	Description	Voice Cloning Method
XTTS	Coqui X-TTS multilingual model	Zero-shot from ~ 10s reference audio
YourTTS	Multi-speaker TTS (Casanova et al., 2022)	Cross-lingual voice conversion

5.3 Semantic Tampering Pipeline



Semantic Tampering Pipeline Diagram

Figure 5.2: Semantic tampering pipeline showing NLP-guided phoneme-boundary editing.

For semantic tampering, edits are made at phoneme boundaries to minimize artifacts:

Tools Used: - Whisper ASR: Word-level transcription with timestamps - Montreal Forced Aligner (MFA): Phoneme-level alignment - spaCy NLP: Part-of-speech tagging and dependency parsing - librosa: Prosody analysis (F0, energy, duration)

Tampering Operations:

Operation	Description	Example
Deletion	Remove adjectives/adverbs	“always on” “on”
Insertion	Add negations	“I agree” “I not agree”
Substitution	Swap quantifiers	“all done” “some done”

6. Training Methodology

6.1 AASIST Training Configuration

Parameter	Value
Batch Size	24
Epochs	100
Learning Rate	1e-4
LR Min	5e-6
Optimizer	Adam
Scheduler	Cosine Annealing
Weight Decay	1e-4
Loss	Weighted CE [0.1, 0.9]
Input Samples	64,600
Sample Rate	16 kHz
First Conv	128
GAT Dims	[64, 32]
Pool Ratios	[0.5, 0.7, 0.5, 0.5]
Hardware	NVIDIA RTX 4080 (16GB)
Training Time	~ 14 hours

Training Progress: - Epoch 13: Best EER 3.03%, tDCF 0.0796 - Epoch 21: Best EER 2.20%, tDCF 0.0615 - Epoch 44: Best EER 1.32%, tDCF 0.0384 (Best tDCF) - Epoch 97: Best EER 0.35% - Final (SWA): EER 1.73%, tDCF 0.0559

6.2 XLS-R + SLS Training Configuration

Parameter	Value
Batch Size	5
Epochs	4 (best: 2)
Learning Rate	1e-6
Optimizer	Adam
Loss	Weighted CE [0.1, 0.9]
Early Stopping	Patience 1
Input Samples	64,600
Sample Rate	16 kHz
Pre-trained Model	XLS-R 300M
Pre-training Hours	436,000
Pre-training Langs	128
Data Augmentation	RawBoost Algo 3
Total Parameters	340M
Hardware	NVIDIA RTX 4080 (16GB)
Training Time	~ 70 minutes

Critical Finding - Overfitting Analysis:

Metric	4 Epochs	100 Epochs
EER	2.97%	44.24%
Training Loss	0.00165	0.00001
Status	Optimal	Overfitting


Lower training loss does NOT indicate better test performance. Early stopping is essential.

7. Experimental Results

7.1 Quantitative Results Summary

Model	Dataset	EER (%)	min t-DCF	ROC AUC	Status
AASIST (Pretrained)	ASVspoof 2019 LA	0.83	0.0275	N/A	Benchmark
AASIST (Trained)	ASVspoof 2019 LA	~ 0.83	~ 0.03	N/A	Reproduced
AASIST (Pretrained)	ASVspoof 2021 LA	50.07	0.7704	0.51	Failed
AASIST (Trained)	ASVspoof 2021 LA	48.27	0.7212	0.55	Failed
XLS-R + SLS (4 ep)	ASVspoof 2019 LA	0.26	N/A	0.999	BEST
XLS-R + SLS (4 ep)	ASVspoof 2021 LA	2.97	0.2674	0.996	BEST
XLS-R + SLS (Paper)	ASVspoof 2021 LA	2.87	N/A	N/A	Target
XLS-R + SLS (100 ep)	ASVspoof 2021 LA	44.24	0.6154	0.58	O�erfitting

7.2 ROC Curve Comparison




ROC Curve Comparison

Figure 7.1: ROC Curve Comparison on ASVspoof 2021 LA. XLS-R + SLS achieves AUC= 0.996, dramatically outperforming AASIST variants (AUC ~ 0.55-0.79).

Key Observations:

- XLS-R + SLS curve hugs the top-left corner (near-perfect classification)
- AASIST curves are close to diagonal (random classifier)
- Self-supervised pre-training provides significant discriminative power

7.3 DET Curve Comparison



DET Curve Comparison

Figure 7.2: Detection Error Tradeoff (DET) curves. Lower-left is better. XLS-R achieves 2.97% EER vs AASIST's 48%.

Key Observations:

- XLS-R + SLS: EER = 2.97% (intersection with diagonal)
- AASIST models: EER > 48% (near-random performance)
- DET curves better visualize performance at low error rates

7.4 Score Distribution Analysis




Figure 7.3: Score distributions for bonafide and spoof samples. Clear separation indicates better discrimination.

Score Statistics (XLS-R + SLS):

Class	Mean	Std	Min	Max	Median
Bonafide	-0.29	1.48	-14.92	0.00	0.00
Spoof	-13.52	3.57	-19.20	0.00	-14.39

Key Insight: XLS-R produces well-separated score distributions with minimal overlap, enabling reliable threshold selection.

7.5 Per-Codec EER Analysis (ASVspoof 2021 LA)





Figure 7.4: Per-codec EER comparison showing XLS-R's superior codec robustness.

Codec	XLS-R EER (%)	AASIST EER (%)	XLS-R Advantage
none (clean)	0.36	22.40	62x better
g722	0.72	24.98	35x better
ulaw	0.81	22.48	28x better
alaw	0.98	22.40	23x better
pstr	2.17	45.69	21x better
opus	2.42	36.17	15x better
gsm	2.59	34.16	13x better

Key Finding: XLS-R maintains < 3% EER across ALL codec conditions while AASIST degrades to near-random performance (22-46% EER) on codec-distorted audio.

7.6 Per-Attack EER Analysis



Per-Attack EER Analysis

Figure 7.5: Per-attack EER for XLS-R + SLS on ASVspoof 2021 LA (attacks A07-A19).

Attack	EER (%)	Difficulty	Description
A13	0.02	Easy	Neural network TTS
A09	0.27	Easy	Vocoder-based
A07	0.67	Easy	TTS system
A08	0.84	Easy	TTS system
A14	0.81	Easy	Neural vocoder
A16	1.26	Easy	Voice conversion
A15	1.49	Easy	Voice conversion
A17	1.74	Easy	Hybrid system
A19	1.78	Easy	End-to-end
A12	2.44	Moderate	TTS + vocoder
A18	2.77	Moderate	Hybrid system
A10	4.41	Challenging	Neural TTS
A11	5.22	Challenging	Neural TTS

Key Finding: Most attacks (11/13) achieve < 3% EER. Only A10 and A11 (neural TTS variants) present moderate challenge.

7.7 Confusion Matrix





Figure 7.6: Confusion matrix for XLS-R + SLS at EER threshold.

Confusion Matrix at EER Threshold:

	Predicted Spoof	Predicted Bonafide
Actual Spoof	129,399 (TP)	3,961 (FN)
Actual Bonafide	440 (FP)	14,376 (TN)

Performance Metrics: - Accuracy at EER: 97.03% - True Positive Rate (Spoof Detection): 97.03% - True Negative Rate (Bonafide Acceptance): 97.03%

7.8 Trans-Splicing Detection Results




Tampering Detection Rates

Figure 7.7: Detection rates for trans-spliced audio across different TTS systems.

Model	XTTS-Clean	XTTS-Unclean	YourTTS-Clean	YourTTS-Unclean	Overall
XLS-R + SLS	91.11%	91.73%	99.81%	100.00%	95.45%
AASIST (Pretrained)	13.44%	11.42%	88.62%	59.95%	42.96%
AASIST (Trained)	8.70%	9.25%	91.98%	58.12%	41.72%

7.9 Model Comparison on Tampering



Model Comparison on Tampering

Figure 7.8: Overall model comparison on trans-splicing detection.

Key Observations:

- XLS-R achieves 95.45% overall detection rate
- AASIST performs poorly on XTTS (~10% detection)
- YourTTS easier to detect for all models

7.10 TTS System Comparison




Figure 7.9: Detection rates by TTS system showing XTTS is harder to detect.

Key Finding: XTTS-generated audio is significantly harder to detect than YourTTS:

- XLS-R: 91% (XTTS) vs 100% (YourTTS)

7.11 Semantic Tampering Results

Model	EER (%)	Notes
XLS-R + SLS	44.44	Domain mismatch
AASIST (Pretrained)	48.00	Near-random
AASIST (Trained)	49.15	Near-random

Limitation: All models struggle with semantic tampering due to: 1. Domain mismatch (different source audio distribution) 2. Subtle nature of phoneme-boundary edits 3. Small dataset size (50 samples)

8. Key Findings

8.1 Scientific Contributions

1. Self-supervised pre-training is critical for robustness

XLS-R maintains < 3% EER across all codec conditions

AASIST fails on codec-distorted audio (22-46% EER)

Pre-training on 436K hours provides codec-invariant representations

2. TTS-specific detection varies by system

XTTS harder to detect than YourTTS for all models

XLS-R detects 91% XTTS vs AASIST's 10%

Modern zero-shot TTS poses greater challenge

3. Overfitting in XLS-R fine-tuning

Early stopping at 2-4 epochs is critical

Extended training (100 epochs) causes catastrophic forgetting

Lower training loss better generalization

4. Codec sensitivity is AASIST's failure mode

Sinc filterbank learns spectral artifacts destroyed by lossy compression

Even clean audio subset shows degraded performance (~ 22% EER)

8.2 Limitations

1. Semantic tampering dataset shows domain mismatch (different source audio)
2. Trans-splicing dataset lacks original bona fide samples for EER calculation
3. Results specific to TTS systems used (XTTS, YourTTS)
4. Models tested only on ASVspoof datasets

8.3 Recommendations

For Researchers:

- Test models on diverse tampering techniques beyond ASVspoof
- Investigate why XTTS is harder to detect than YourTTS
- Develop dedicated audio forensics models for semantic tampering

For Production Systems:

- Use XLS-R + SLS for audio deepfake detection
- Apply early stopping (2-4 epochs) during fine-tuning
- Monitor for novel TTS systems and attack vectors

Part II: Reproducibility Guide

9. Environment Setup

9.1 Hardware Requirements

Component	Minimum	Recommended
GPU	8GB VRAM	16GB VRAM (RTX 3080/4080)
RAM	16GB	32GB
Storage	50GB	100GB (including datasets)
CPU	4 cores	8+ cores

9.2 Software Installation

```
# Create conda environment
conda create -n deepfake_detection python=3.8
conda activate deepfake_detection

# Install PyTorch with CUDA support
pip install torch==1.12.1 torchvision==0.13.1 torchaudio==0.12.1 --extra-index-url https://download.pytorch.org/wheel/cu116

# Install dependencies
pip install soundfile librosa numpy pandas matplotlib scipy scikit-learn
pip install gradio>=3.34.0

# For AAST
pip install torchcontrib

# For XLS-R (fairseq)
cd xlsr_sls/SLSforASVspoof-2021-DF
pip install -e fairseq-a54021305d6b3c4c5959ac9395135f63202db8f1/
```

10. Dataset Preparation

10.1 ASVspoof 2019 LA

```
# Download from Edinburgh DataShare
# URL: https://datashare.ed.ac.uk/handle/10283/3336

# Expected directory structure after extraction:
data/asvspoof/asvspoof2019/LA/
    ASVspoof2019_LA_train/
        flac/
    ASVspoof2019_LA_dev/
        flac/
    ASVspoof2019_LA_eval/
        flac/
    ASVspoof2019_LA_cm_protocols/
        ASVspoof2019_LA_cm_train.trn.txt
        ASVspoof2019_LA_cm_dev.trl.txt
        ASVspoof2019_LA_cm_eval.trl.txt
```

10.2 ASVspoof 2021 LA

```
# Download from Zenodo
# Audio: https://zenodo.org/record/4837263
# Keys: https://www.asvspoof.org/index2021.html

# Expected directory structure:
data/asvspoof/asvspoof2021/
    ASVspoof2021_LA_eval/
        flac/
        keys/
        LA/
        CM/
            trial_metadata.txt
```

11. Training Procedures

11.1 Training AASIST

```
cd aasi_st

# Edit config file to set your dataset path
# In config/AASI ST.conf, update:
# "database_path": "/path/to/your/data/asvspoof/asvspoof2019/LA/"

# Start training
python main.py --config config/AASI ST.conf

# Training will create:
# exp_result/LA_AASI ST_ep100_bs24/
#     weights/          # Model checkpoints
#     metrics/          # Evaluation scores
#     metric_log.txt    # Training log
```

Monitoring Training: - Check `metric_log.txt` for epoch-wise metrics - Best model saved when dev EER improves - Training completes in ~ 14 hours on RTX 4080

11.2 Training XLS-R + SLS

```
cd xls_r_sls/SLSforASVspoof-2021-DF

# Download XLS-R pretrained weights (if not present)
# xlsr2_300m.pt should be in the directory

# Edit train_LA.sh to set your paths

# Start training
bash train_LA.sh

# Training will create:
# models/model_LA_WCE_50_5_1e-06_*/
#     epoch_*.pth      # Checkpoints
#     training_output_LA.log
```

Important: Stop training after 2-4 epochs (best EER achieved at epoch 2).

12. Evaluation Procedures

12.1 Evaluating AASIST

```
cd aasi st

# Using pretrained model
python main.py --eval --config config/AASI ST.conf

# Using your trained model
# Edit config to set model_path to your checkpoint
python main.py --eval --config config/AASI ST.conf
```

12.2 Evaluating XLS-R + SLS

```
cd xls_r_sls/SLSforASVspoof-2021-DF

# Evaluate on ASVspoof 2021 LA
python eval_LA.py \
    --model_path best_model_4epochs_2.97EER.pth \
    --data_path /path/to/ASVspoof2021_LA_eval/flac \
    --protocol /path/to/keys/LA/CM/trial_metadata.txt

# Output: scores_LA_epoch*.txt
```

12.3 Evaluating on Tampering Datasets

```
cd tampered_evaluation


# Evaluate XLS-R on Trans-Slicing
python eval_tampered.py --model xlsr --dataset trans_slicing

# Evaluate AASIST on Trans-Slicing
python eval_tampered.py --model aasi st --model_variant pretrained --dataset trans_slicing

# Evaluate all combinations
python eval_tampered.py --model all --dataset all
```

13. GUI Applications

13.1 AASIST Detection Interface



AASIST GUI

Figure 13.1: AASIST single-file detection interface with waveform visualization.


Launch Command:

```
cd aasi st
python gradio_app.py --port 7860
# Access at http://127.0.0.1:7860
```

Features:

- Upload audio or record from microphone
- Select model checkpoint
- View prediction scores and visualizations
- Waveform and spectrogram display

13.2 AASIST Multi-Tab Interface



AASIST Multi-Tab GUI


Figure 13.2: AASIST multi-tab interface with batch processing and model comparison capabilities.

Launch Command:

```
cd aasist
python gradio_app_multitab.py --port 7860
# Access at http://127.0.0.1:7860
```

Tabs: 1. Single Detection: Analyze individual audio files 2. Batch Processing: Process multiple files with CSV export 3. Model Comparison: Compare predictions from different models 4. Training Monitor: View training progress from log files 5. Dataset Explorer: Browse ASVspoof datasets

13.3 XLS-R + SLS Interface



XLS-R GUI

Figure 13.3: XLS-R + SLS detection interface with color-coded results.

Launch Command:

```
cd xl s_r_sI s/SLSF or ASVsspoof - 2021- DF  
python gradi o_app. py --por t 7861  
# Access at http://127. 0. 0. 1: 7861
```

Features: - Upload audio or record from microphone - XLS-R model selection - Color-coded results (green= real, red= fake) - Waveform visualization

13.4 Running Both GUIs Simultaneously

```
# Terminal 1: AASI ST (port 7860)
cd aasi_st && python gradi o_app_muli tab. py

# Terminal 2: XLS-R (port 7861)
cd xls_r_sls/SLSf or ASVspoof - 2021- DF && python gradi o_app. py
```

14. Troubleshooting

14.1 Common Issues

Issue	Solution
CUDA out of memory	Reduce batch_size in config
Model not loading	Check checkpoint path is correct
Audio format error	Convert to WAV/FLAC, 16kHz mono
Fairseq import error	Reinstall fairseq from source
Gradio version error	Update to gradio>= 3.34.0

14.2 GPU Memory Requirements

Model	Batch Size	GPU Memory
AASIST	24	~ 8GB
AASIST	12	~ 4GB
XLS-R	5	~ 12GB
XLS-R	2	~ 6GB

14.3 Known Limitations

1. XLS-R requires fairseq: Must install specific fairseq version from source
 2. AASIST codec sensitivity: Poor performance on codec-distorted audio
 3. GPU required: CPU training not supported
-

Part III: References

15. Scientific Citations

Models

```

@article{jung2022aasi_st,
    title={AASI ST: Audio anti-spoofing using integrated spectro-temporal graph attention networks},
    author={Jung, Jee-weon and Heo, Hee-Soo and Tak, Hemata and Shim, Hye-jin and Chung, Joon Son and Lee, Bong-Jin and Yu, Ha-Jin and Evans, Nicholas},
    journal={IEEE/ACM Transactions on Audio, Speech, and Language Processing},
    volume={30},
    pages={1592--1603},
    year={2022}
}

@proceedings{zhang2024audio,
    title={Audio Deepfake Detection with Self-supervised XLS-R and SLS classifier},
    author={Zhang, Qishan and Wen, Shuangbing and Hu, Tao},
    booktitle={Proceedings of the 32nd ACM International Conference on Multimedia},
    pages={10873--10877},
    year={2024}
}

@article{babu2021xlsr,
    title={XLS-R: Self-supervised cross-lingual speech representation learning at scale},
    author={Babu, Arun and Wang, Changhan and Tjandra, Andros and others},
    journal={arXiv preprint arXiv: 2111.09296},
    year={2021}
}

```

Datasets

```

@article{wang2020asvspoof,
  title={ASVspoof 2019: A Large-scale public database of synthesized, converted and replayed speech},
  author={Wang, Xin and Yamagishi, Junichi and Todisco, Massimiliano and others},
  journal={Computer Speech \& Language},
  volume={64},
  pages={101114},
  year={2020}
}

@article{yamagishi2022asvspoof,
  title={ASVspoof 2021: Towards spoofed and deepfake speech detection in the wild},
  author={Yamagishi, Junichi and Wang, Xin and Todisco, Massimiliano and others},
  journal={IEEE/ACM Transactions on Audio, Speech, and Language Processing},
  volume={30},
  pages={2507--2522},
  year={2022}
}

```

Evaluation Metrics

```

@proceedings{kinnunen2018t,
  title={t-DCF: a detection cost function for the tandem assessment of spoofing countermeasures and automatic speaker verification},
  author={Kinnunen, Tomi and Lee, Kong Alk and Delgado, Hector and Evans, Nicholas and Todisco, Massimiliano and Sahidullah, Md and Yamagishi, Junichi and Reynolds, Douglas A},
  booktitle={Proc. Odyssey},
  pages={312--319},
  year={2018}
}

```

TTS Systems

```
@inproceedings{casanova2022your tts,
  title={Your TTS: Towards Zero-Shot Multi-Speaker TTS and Zero-Shot Voice Conversion for Everyone},
  author={Casanova, Edresson and Weber, Julian and Shulby, Christopher and Junior, Arnaldo Cando and Gólgógi, Eren and Ponti, Macir Antonelli},
  booktitle={International Conference on Machine Learning},
  pages={2709- - 2720},
  year={2022}
}
```

Repository Structure

```
deepfake_models/
  aasi_st/
    config/
    models/
      weights/
      exp_result/
      results/
      main.py
      gradio_app.py
      gradio_app_multitab.py
  xls_r_sls/
    SLSforASVspoof - 2021-DF/
      model.py
      figures/
      train_LA.sh
      gradio_app.py
      best_model_*.pth
  tampered_evaluation/
    trans_slicing/
      semantic/
      eval_tampered.py
  figures/
    gui_*.png
    trans_slicing_digram.png # Pipeline diagram
    semantic_tampering_digram.png
    tampering_*.png
  project_report/
  data/
```

AASI ST model implementation
Configuration files
Model architectures
Pretrained weights
Training outputs
Evaluation results & plots
Training script
Simple GUI
Multi-tab GUI
XLS-R + SLS implementation
Model architecture
Evaluation visualizations
Training script
GUI application
Best checkpoint
Tampering evaluation
Trans-slicing dataset
Semantic tampering dataset
Evaluation script
Project visualizations
GUI screenshots
Pipeline diagram
Semantic tampering diagram
Tampering evaluation plots
Generated reports
Datasets (not in repo)

Code Availability

Project Repository

Main Repository: https://github.com/Yash-Sukhdeve/deepfake_models

This repository contains: - AASIST model implementation and pretrained weights - XLS-R + SLS model implementation - Tampering evaluation scripts - GUI applications for both models - All configuration files and training scripts

Tampering Dataset Repositories

Repository	Description	Files	URL
Audio-Tampering	Trans-splicing dataset with TTS-generated word insertions using XTTS and YourTTS	1,932	https://github.com/AVHBAC/Audio-Tampering
Tampered_Deepfake	Semantic tampering dataset with NLP-guided deletions, insertions, and substitutions at phoneme boundaries	50	https://github.com/Yash-Sukhdeve/Tampered_Deepfake

External Model Repositories

Repository	Description	URL
XLS-R + SLS	Original implementation (Zhang et al., ACM MM 2024)	https://github.com/QiShanZhang/SLSforASVspoof-2021-DF
ASVspoof 2021 Baseline	Official ASVspoof 2021 baseline with RawNet2	https://github.com/asvspoof-challenge/2021
RawGAT-ST	Spectro-temporal graph attention networks for anti-spoofing	https://github.com/eurecom-asp/RawGAT-ST-antispoofing
Fairseq	Facebook AI Research toolkit (XLS-R backbone)	https://github.com/pytorch/fairseq

Tools and Libraries

Tool	Purpose	URL
Coqui TTS (XTTS)	Zero-shot multilingual TTS for trans-splicing	https://github.com/coqui-ai/TTS
YourTTS	Multi-speaker TTS for voice cloning	https://github.com/Edresson/YourTTS
Montreal Aligner	Forced Phoneme-level alignment for semantic tampering	https://github.com/MontrealCorpusTools/Montreal-Forced-Aligner
Whisper	Word-level transcription and alignment	https://github.com/openai/whisper

Conclusions

1. Self-supervised pre-training dramatically improves audio deepfake detection performance across diverse evaluation conditions.
2. XLS-R + SLS achieves 2.97% EER on ASVspoof 2021 LA, successfully reproducing the ACM MM 2024 paper results (target: 2.87% EER).
3. AASIST fails to generalize to ASVspoof 2021 due to codec sensitivity, despite achieving benchmark performance (0.83% EER) on ASVspoof 2019 LA.
4. The 87% of ASVspoof 2021 samples with codec distortion expose models trained on clean audio - highlighting the importance of realistic evaluation.
5. Early stopping is critical: Extended training leads to severe overfitting (2.97% → 44.24% EER for XLS-R + SLS after 100 epochs).
6. XLS-R + SLS detects 95.45% of trans-spliced audio, significantly outperforming AASIST (41.72%).

— End of Report - November 2025